

```
In [1]: ❸ import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: ❸ import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
In [3]: ❸ train_data_generator = ImageDataGenerator(rescale = 1./255,
                                                    shear_range = 0.2,
                                                    zoom_range = 0.2,
                                                    horizontal_flip = True)
```

```
In [4]: ❸ train_set = train_data_generator.flow_from_directory('train',
                                                            target_size = (32, 32),
                                                            batch_size = 16,
                                                            class_mode = 'binary')
```

Found 40 images belonging to 2 classes.

```
In [5]: ❸ train_set_filenames[0]
```

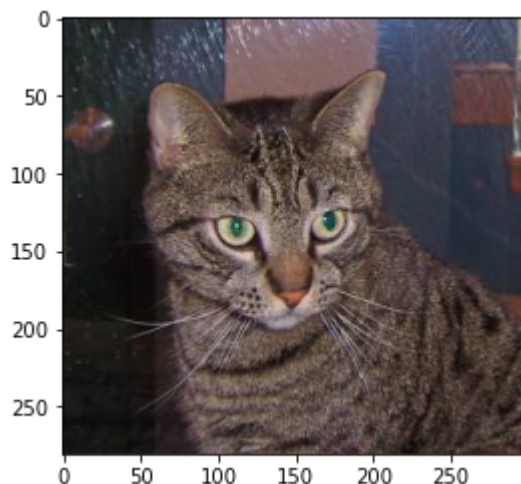
Out[5]: 'cats\\1.jpg'

```
In [6]: ❸ img = tf.keras.preprocessing.image.load_img('train\\cats\\1.jpg')
```

```
In [7]: ❸ import matplotlib.pyplot as plt
```

```
In [8]: ❸ plt.imshow(img)
```

Out[8]: <matplotlib.image.AxesImage at 0x1ce95868c40>



```
In [9]: train_set.class_indices
```

```
Out[9]: {'cats': 0, 'dogs': 1}
```

```
In [10]: test_data_generator = ImageDataGenerator(rescale = 1./255)
```

```
In [11]: test_set = test_data_generator.flow_from_directory('test',
                                                         target_size = (32,32),
                                                         batch_size = 16,
                                                         class_mode = 'binary')
```

Found 20 images belonging to 2 classes.

```
In [57]: from tensorflow.keras.models import Sequential
```

```
In [58]: from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout
```

```
In [59]: model = Sequential()
```

Convolutional layer 1 with 32 filters of kernel size[5,5]

```
In [60]: model.add(Conv2D(filters = 32,
                          kernel_size = [5,5],
                          activation = 'relu',
                          padding = 'valid',
                          input_shape = [32,32,3]))
```

Pooling layer 1 with pool size[2,2] and stride 2

```
In [61]: model.add(MaxPool2D(pool_size = [2,2],
                              strides = 2))
```

Convolutional layer 2 with 64 filters of kernel size[5,5]

```
In [62]: model.add(Conv2D(filters = 64,
                          kernel_size = [5,5],
                          activation = 'relu',
                          padding = 'valid'))
```

Pooling layer 2 with pool size[2,2] and stride 2

```
In [63]: model.add(MaxPool2D(pool_size = [2,2],  
                             strides = 2))
```

```
In [64]: model.add(Flatten())
```

Dense layer whose output size is fixed in the hyper parameter: fc_size=32

```
In [65]: model.add(Dense(units = 32,  
                          activation = 'relu'))
```

Dropout layer with dropout probability 0.4

```
In [66]: model.add(Dropout(0.40))
```

Predict the class by doing a sigmoid on the output of the dropout layers.

```
In [67]: model.add(Dense(units = 1,  
                          activation = 'sigmoid'))
```

In [68]: `model.summary()`

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 28, 28, 32)	2432
max_pooling2d_4 (MaxPooling 2D)	(None, 14, 14, 32)	0
conv2d_5 (Conv2D)	(None, 10, 10, 64)	51264
max_pooling2d_5 (MaxPooling 2D)	(None, 5, 5, 64)	0
flatten_2 (Flatten)	(None, 1600)	0
dense_4 (Dense)	(None, 32)	51232
dropout_2 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 1)	33
=====		
Total params: 104,961		
Trainable params: 104,961		
Non-trainable params: 0		
=====		

For the training step, define the loss function and minimize it

In [69]: `model.compile(optimizer = 'adam',
loss = 'binary_crossentropy',
metrics = ['accuracy'])`

In [25]: `model.fit(train_set, validation_data = test_set, epochs = 100)`

```
3/3 [=====] - 0s 160ms/step - loss: 0.0557 - ac
curacy: 1.0000 - val_loss: 1.0487 - val_accuracy: 0.7000
Epoch 82/100
3/3 [=====] - 0s 161ms/step - loss: 0.0364 - ac
curacy: 1.0000 - val_loss: 1.0387 - val_accuracy: 0.6500
Epoch 83/100
3/3 [=====] - 0s 142ms/step - loss: 0.1005 - ac
curacy: 0.9500 - val_loss: 0.9699 - val_accuracy: 0.6500
Epoch 84/100
3/3 [=====] - 0s 163ms/step - loss: 0.0535 - ac
curacy: 1.0000 - val_loss: 0.9231 - val_accuracy: 0.6500
Epoch 85/100
3/3 [=====] - 0s 161ms/step - loss: 0.1307 - ac
curacy: 0.9750 - val_loss: 0.9241 - val_accuracy: 0.6000
Epoch 86/100
3/3 [=====] - 0s 140ms/step - loss: 0.1090 - ac
curacy: 0.9500 - val_loss: 0.8887 - val_accuracy: 0.6000
Epoch 87/100
3/3 [=====] - 0s 173ms/step - loss: 0.0579 - ac
curacy: 1.0000 - val_loss: 0.8754 - val_accuracy: 0.6500
```

In [48]: `model.fit(train_set, validation_data = test_set, epochs = 200)`

```
Epoch 195/200
3/3 [=====] - 0s 116ms/step - loss: 0.0544 - ac
curacy: 0.9750 - val_loss: 1.2127 - val_accuracy: 0.7500
Epoch 196/200
3/3 [=====] - 0s 112ms/step - loss: 0.0332 - ac
curacy: 1.0000 - val_loss: 1.1159 - val_accuracy: 0.7000
Epoch 197/200
3/3 [=====] - 0s 109ms/step - loss: 0.0254 - ac
curacy: 1.0000 - val_loss: 1.1362 - val_accuracy: 0.6500
Epoch 198/200
3/3 [=====] - 0s 95ms/step - loss: 0.0081 - acc
uracy: 1.0000 - val_loss: 1.1631 - val_accuracy: 0.6500
Epoch 199/200
3/3 [=====] - 0s 141ms/step - loss: 0.0309 - ac
curacy: 0.9750 - val_loss: 1.1101 - val_accuracy: 0.7500
Epoch 200/200
3/3 [=====] - 0s 133ms/step - loss: 0.0077 - ac
curacy: 1.0000 - val_loss: 1.1017 - val_accuracy: 0.7000
```

Out[48]: `<keras.callbacks.History at 0x1ce98dbfc40>`

In [70]: `model.fit(train_set, validation_data = test_set, epochs = 300)`

```
3/3 [=====] - 0s 119ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 1.0800 - val_accuracy: 0.8000
Epoch 281/300
3/3 [=====] - 0s 86ms/step - loss: 0.0214 - accuracy: 1.0000 - val_loss: 1.0913 - val_accuracy: 0.8000
Epoch 282/300
3/3 [=====] - 0s 99ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 1.0980 - val_accuracy: 0.8000
Epoch 283/300
3/3 [=====] - 0s 77ms/step - loss: 0.0044 - accuracy: 1.0000 - val_loss: 1.0947 - val_accuracy: 0.8000
Epoch 284/300
3/3 [=====] - 0s 78ms/step - loss: 0.0187 - accuracy: 0.9750 - val_loss: 1.0835 - val_accuracy: 0.8000
Epoch 285/300
3/3 [=====] - 0s 80ms/step - loss: 0.0099 - accuracy: 1.0000 - val_loss: 1.0744 - val_accuracy: 0.8000
Epoch 286/300
3/3 [=====] - 0s 85ms/step - loss: 0.0166 - accuracy: 1.0000 - val_loss: 1.0358 - val_accuracy: 0.8500
```

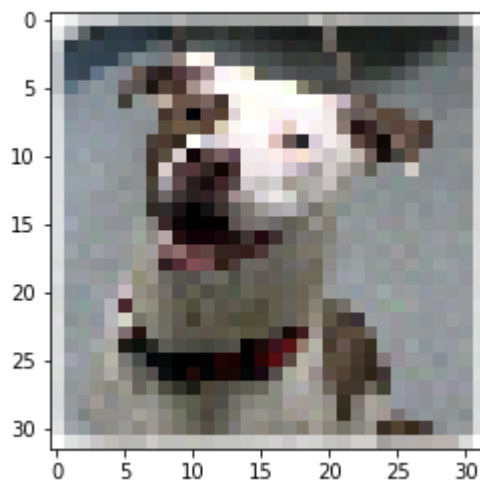
In [71]: `import numpy as np`

In [72]: `from tensorflow.keras.preprocessing import image`

In [73]: `test_img = image.load_img('test/dogs/101.jpg',target_size = (32,32,3))`

In [74]: `plt.imshow(test_img)`

Out[74]: `<matplotlib.image.AxesImage at 0x1ce9c2c7790>`



In [75]: `test_img = image.img_to_array(test_img)`
`test_img = np.expand_dims(test_img, axis=0)`
`test_img = test_img/255`

```
In [76]: ▶ result = model.predict(test_img)
```

```
In [77]: ▶ result
```

```
Out[77]: array([[1.]], dtype=float32)
```

```
In [78]: ▶ if np.round(result[0][0]) == 1:  
           prediction = 'dog'  
       else:  
           prediction = 'cat'  
       print(prediction)
```

dog